

iscafr6fj

February 22, 2025

```
[1]: import numpy as np
import pandas as pd
```

Get the data

```
[2]: col_names = ['sepal_length', 'sepal_width', 'petal_length', 'petal_width',
    ↪ 'type']
data = pd.read_csv("Iris.csv", skiprows=1, header=None, names=col_names)
data.head(10)
```

```
[2]:
```

	sepal_length	sepal_width	petal_length	petal_width	type
1	5.1	3.5	1.4	0.2	Iris-setosa
2	4.9	3.0	1.4	0.2	Iris-setosa
3	4.7	3.2	1.3	0.2	Iris-setosa
4	4.6	3.1	1.5	0.2	Iris-setosa
5	5.0	3.6	1.4	0.2	Iris-setosa
6	5.4	3.9	1.7	0.4	Iris-setosa
7	4.6	3.4	1.4	0.3	Iris-setosa
8	5.0	3.4	1.5	0.2	Iris-setosa
9	4.4	2.9	1.4	0.2	Iris-setosa
10	4.9	3.1	1.5	0.1	Iris-setosa

Node classes

```
[3]: class Node():
    def __init__(self, feature_index=None, threshold=None, left=None,
    ↪ right=None, info_gain=None, value=None):
        ''' constructor '''

        # for decision node
        self.feature_index = feature_index
        self.threshold = threshold
        self.left = left
        self.right = right
        self.info_gain = info_gain

        # for leaf node
        self.value = value
```

Tree Class

```
[4]: class DecisionTreeClassifier():
    def __init__(self, min_samples_split=2, max_depth=2):
        ''' constructor '''

        # initialize the root of the tree
        self.root = None

        # stopping conditions
        self.min_samples_split = min_samples_split
        self.max_depth = max_depth

    def build_tree(self, dataset, curr_depth=0):
        ''' recursive function to build the tree '''

        X, Y = dataset[:, :-1], dataset[:, -1]
        num_samples, num_features = np.shape(X)

        # split until stopping conditions are met
        if num_samples >= self.min_samples_split and curr_depth <= self.max_depth:
            # find the best split
            best_split = self.get_best_split(dataset, num_samples, num_features)
            # check if information gain is positive
            if best_split["info_gain"] > 0:
                # recur left
                left_subtree = self.build_tree(best_split["dataset_left"],
                ↪ curr_depth+1)
                # recur right
                right_subtree = self.build_tree(best_split["dataset_right"],
                ↪ curr_depth+1)
                # return decision node
                return Node(best_split["feature_index"],
                ↪ best_split["threshold"],
                           left_subtree, right_subtree,
                ↪ best_split["info_gain"])

            # compute leaf node
            leaf_value = self.calculate_leaf_value(Y)
            # return leaf node
            return Node(value=leaf_value)

    def get_best_split(self, dataset, num_samples, num_features):
        ''' function to find the best split '''

        # dictionary to store the best split
        best_split = {}
```

```

max_info_gain = -float("inf")

# loop over all the features
for feature_index in range(num_features):
    feature_values = dataset[:, feature_index]
    possible_thresholds = np.unique(feature_values)
    # loop over all the feature values present in the data
    for threshold in possible_thresholds:
        # get current split
        dataset_left, dataset_right = self.split(dataset,
↪feature_index, threshold)
        # check if childs are not null
        if len(dataset_left)>0 and len(dataset_right)>0:
            y, left_y, right_y = dataset[:, -1], dataset_left[:, -1],
↪dataset_right[:, -1]
            # compute information gain
            curr_info_gain = self.information_gain(y, left_y, right_y,
↪"gini")

            # update the best split if needed
            if curr_info_gain>max_info_gain:
                best_split["feature_index"] = feature_index
                best_split["threshold"] = threshold
                best_split["dataset_left"] = dataset_left
                best_split["dataset_right"] = dataset_right
                best_split["info_gain"] = curr_info_gain
                max_info_gain = curr_info_gain

# return best split
return best_split

def split(self, dataset, feature_index, threshold):
    ''' function to split the data '''

    dataset_left = np.array([row for row in dataset if
↪row[feature_index]<=threshold])
    dataset_right = np.array([row for row in dataset if
↪row[feature_index]>threshold])
    return dataset_left, dataset_right

def information_gain(self, parent, l_child, r_child, mode="entropy"):
    ''' function to compute information gain '''

    weight_l = len(l_child) / len(parent)
    weight_r = len(r_child) / len(parent)
    if mode=="entropy":
        gain = self.gini_index(parent) - (weight_l*self.gini_index(l_child)
↪+ weight_r*self.gini_index(r_child))

```

```

        else:
            gain = self.entropy(parent) - (weight_l*self.entropy(l_child) +
↪weight_r*self.entropy(r_child))
            return gain

def entropy(self, y):
    ''' function to compute entropy '''

    class_labels = np.unique(y)
    entropy = 0
    for cls in class_labels:
        p_cls = len(y[y == cls]) / len(y)
        entropy += -p_cls * np.log2(p_cls)
    return entropy

def gini_index(self, y):
    ''' function to compute gini index '''

    class_labels = np.unique(y)
    gini = 0
    for cls in class_labels:
        p_cls = len(y[y == cls]) / len(y)
        gini += p_cls**2
    return 1 - gini

def calculate_leaf_value(self, Y):
    ''' function to compute leaf node '''

    Y = list(Y)
    return max(Y, key=Y.count)

def print_tree(self, tree=None, indent=" "):
    ''' function to print the tree '''

    if not tree:
        tree = self.root

    if tree.value is not None:
        print(tree.value)

    else:
        print("X_"+str(tree.feature_index), "<=", tree.threshold, "?", tree.
↪info_gain)
        print("%sleft:" % (indent), end="")
        self.print_tree(tree.left, indent + indent)
        print("%sright:" % (indent), end="")
        self.print_tree(tree.right, indent + indent)

```

```

def fit(self, X, Y):
    ''' function to train the tree '''

    dataset = np.concatenate((X, Y), axis=1)
    self.root = self.build_tree(dataset)

def predict(self, X):
    ''' function to predict new dataset '''

    predictions = [self.make_prediction(x, self.root) for x in X]
    return predictions

def make_prediction(self, x, tree):
    ''' function to predict a single data point '''

    if tree.value!=None: return tree.value
    feature_val = x[tree.feature_index]
    if feature_val<=tree.threshold:
        return self.make_prediction(x, tree.left)
    else:
        return self.make_prediction(x, tree.right)

```

Train Test Split

```

[5]: X = data.iloc[:, :-1].values
     Y = data.iloc[:, -1].values.reshape(-1,1)
     from sklearn.model_selection import train_test_split
     X_train, X_test, Y_train, Y_test = train_test_split(X, Y, test_size=.2,
     ↪random_state=41)

```

Fit the Model

```

[6]: classifier = DecisionTreeClassifier(min_samples_split=3, max_depth=3)
     classifier.fit(X_train,Y_train)
     classifier.print_tree()

```

```

X_2 <= 1.9 ? 0.9264046681474138
left:Iris-setosa
right:X_3 <= 1.5 ? 0.7694993941591152
left:X_2 <= 4.9 ? 0.17556502585750278
left:Iris-versicolor
right:Iris-virginica
right:X_2 <= 5.0 ? 0.1228956258058704
left:X_1 <= 2.8 ? 0.46691718668869925
left:Iris-virginica
right:Iris-versicolor
right:Iris-virginica

```

Test the Model

```
[7]: Y_pred = classifier.predict(X_test)
      from sklearn.metrics import accuracy_score
      accuracy_score(Y_test, Y_pred)
```

```
[7]: 0.9333333333333333
```

##Using scikit learn

```
[8]: # Load libraries
      import pandas as pd
      from sklearn.tree import DecisionTreeClassifier # Import Decision Tree
      ↪Classifier
      from sklearn.model_selection import train_test_split # Import train_test_split
      ↪function
      from sklearn import metrics #Import scikit-learn metrics module for accuracy
      ↪calculation
```

Read Data

```
[9]: col_names = ['sepal_length', 'sepal_width', 'petal_length', 'petal_width',
      ↪'type']
      pima = pd.read_csv("Iris.csv", skiprows=1, header=None, names=col_names)
      pima.head(10)
```

```
[9]:
```

	sepal_length	sepal_width	petal_length	petal_width	type
1	5.1	3.5	1.4	0.2	Iris-setosa
2	4.9	3.0	1.4	0.2	Iris-setosa
3	4.7	3.2	1.3	0.2	Iris-setosa
4	4.6	3.1	1.5	0.2	Iris-setosa
5	5.0	3.6	1.4	0.2	Iris-setosa
6	5.4	3.9	1.7	0.4	Iris-setosa
7	4.6	3.4	1.4	0.3	Iris-setosa
8	5.0	3.4	1.5	0.2	Iris-setosa
9	4.4	2.9	1.4	0.2	Iris-setosa
10	4.9	3.1	1.5	0.1	Iris-setosa

Train Test Split

```
[10]: #split dataset in features and target variable
      feature_cols = ['sepal_length', 'sepal_width', 'petal_length', 'petal_width']
      X = pima[feature_cols] # Features
      y = pima.type # Target variable
```

```
[11]: # Split dataset into training set and test set
      X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3,
      ↪random_state=1) # 70% training and 30% test
```

Fit and predict

```
[12]: # Create Decision Tree classifier object
      clf = DecisionTreeClassifier()
```

```
[13]: # Train Decision Tree Classifier
      clf = clf.fit(X_train,y_train)
```

```
[14]: #Predict the response for test dataset
      y_pred = clf.predict(X_test)
```

Performance Evaluation

```
[15]: # Model Accuracy, how often is the classifier correct?
      print("Accuracy:",metrics.accuracy_score(y_test, y_pred))
```

Accuracy: 0.9555555555555556


```
[16]: import seaborn as sns
      import matplotlib.pyplot as plt
      from sklearn.metrics import confusion_matrix

      # Calculate the confusion matrix
      cm = confusion_matrix(y_test, y_pred)

      # Plot the confusion matrix
      plt.figure(figsize=(8, 6)) # Set the figure size
      sns.heatmap(cm, annot=True, fmt="d", cmap="Blues", xticklabels=pima['type'].
        ↪unique(), yticklabels=pima['type'].unique())

      # Add labels and title
      plt.xlabel('Predicted')
      plt.ylabel('True')
      plt.title('Confusion Matrix')

      # Show the plot
      plt.show()
```



Visualization

```
[17]: pip install graphviz
```

Requirement already satisfied: graphviz in /home/smayan/Desktop/AI-ML-DS/AI-and-ML-Course/.conda/lib/python3.11/site-packages (0.20.3)

[notice] A new release of pip is available: 25.0 -> 25.0.1

[notice] To update, run:

`pip install --upgrade pip`

Note: you may need to restart the kernel to use updated packages.

```
[18]: pip install pydotplus
```

Requirement already satisfied: pydotplus in /home/smayan/Desktop/AI-ML-DS/AI-and-ML-Course/.conda/lib/python3.11/site-packages (2.0.2)

Requirement already satisfied: pyparsing>=2.0.1 in /home/smayan/Desktop/AI-ML-

DS/AI-and-ML-Course/.conda/lib/python3.11/site-packages (from pydotplus) (3.2.1)

[notice] A new release of pip is

available: 25.0 -> 25.0.1

[notice] To update, run:

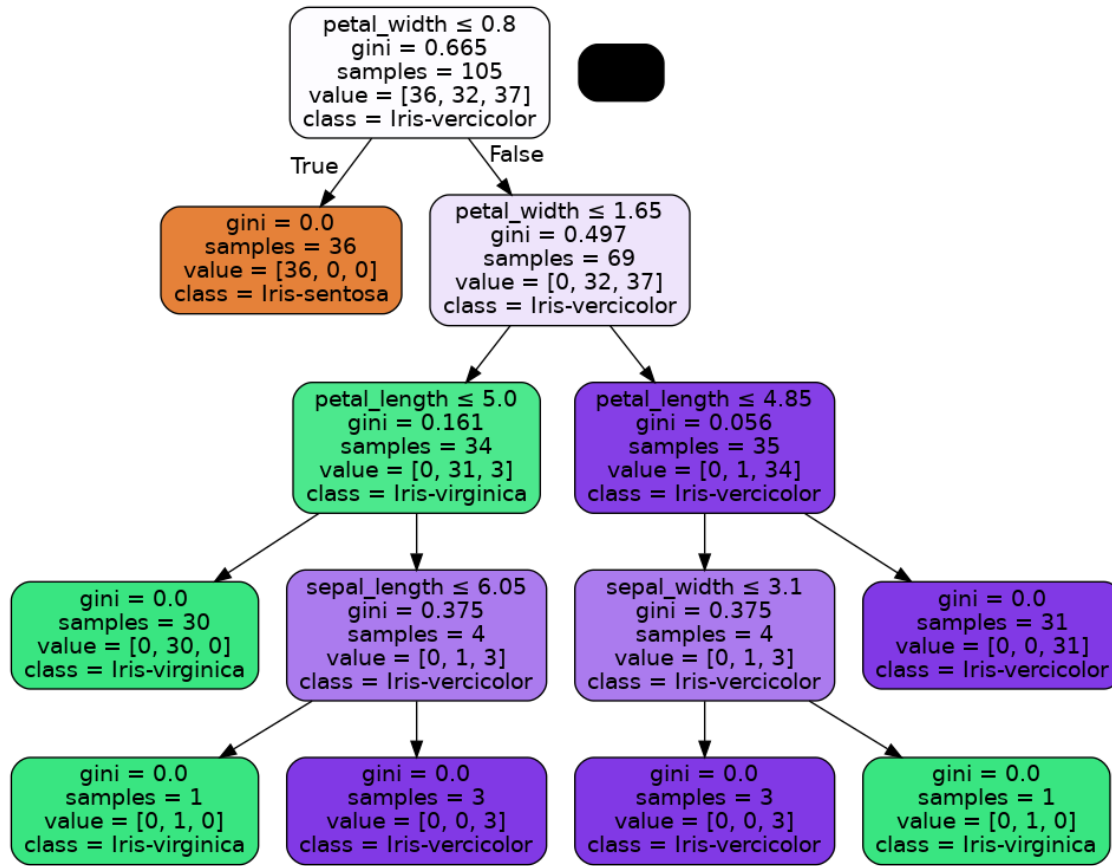
`pip install --upgrade pip`

Note: you may need to restart the kernel to use updated packages.

```
[19]: import six
import sys
sys.modules['sklearn.externals.six'] = six
```

```
[20]: from sklearn.tree import export_graphviz
from IPython.display import Image
import pydotplus
from sklearn.externals.six import StringIO
dot_data = StringIO()
export_graphviz(clf, out_file=dot_data,
                filled=True, rounded=True,
                special_characters=True, feature_names =
↳ feature_cols, class_names=['Iris-sentosa', 'Iris-virginica',
↳ 'Iris-vercicolor'])
graph = pydotplus.graph_from_dot_data(dot_data.getvalue())
graph.write_png('iris.png')
Image(graph.create_png())
```

[20]:



Prediction using Entropy Method

```
[21]: # Create Decision Tree classifier object
clf = DecisionTreeClassifier(criterion="entropy", splitter="random",
    ↪max_depth=3)

# Train Decision Tree Classifier
clf = clf.fit(X_train,y_train)

#Predict the response for test dataset
y_pred = clf.predict(X_test)

# Model Accuracy, how often is the classifier correct?
print("Accuracy:",metrics.accuracy_score(y_test, y_pred))
```

Accuracy: 0.9111111111111111

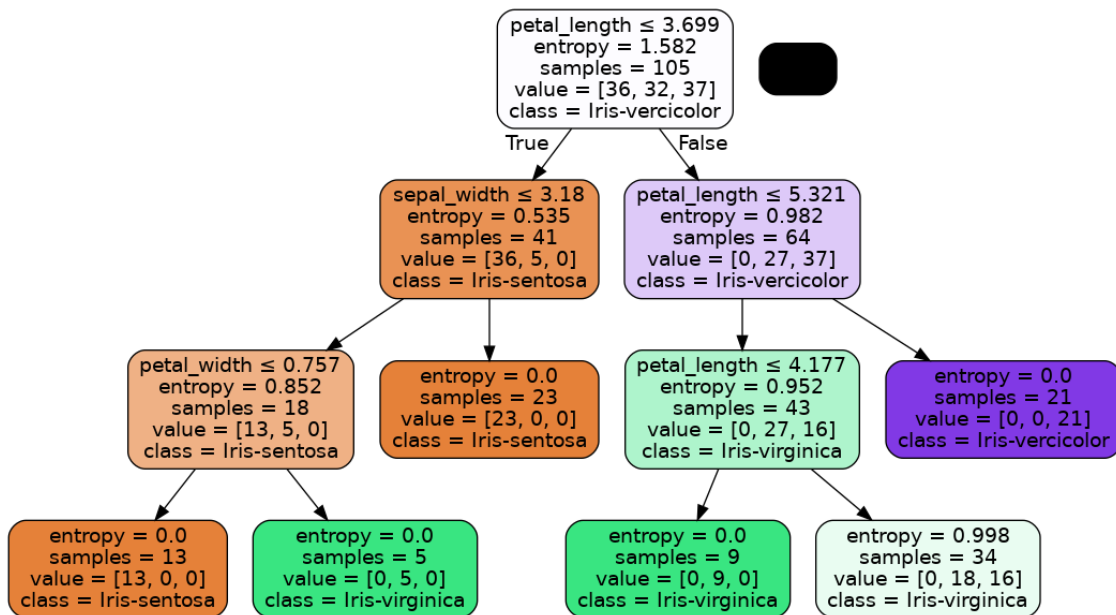
```
[22]: from sklearn.tree import export_graphviz
from IPython.display import Image
import pydotplus
```

```

from sklearn.externals.six import StringIO
dot_data = StringIO()
export_graphviz(clf, out_file=dot_data,
                filled=True, rounded=True,
                special_characters=True, feature_names = _
↳ feature_cols, class_names=['Iris-sentosa', 'Iris-virginica', _
↳ 'Iris-vercicolor'])
graph = pydotplus.graph_from_dot_data(dot_data.getvalue())
graph.write_png('iris.png')
Image(graph.create_png())

```

[22]:



##Overfitting on synthetic data set

[]:

Use the `make_classification()` function to define a binary (two class) classification prediction problem with 10,000 examples (rows) and 20 input features (columns).

```

[23]: # evaluate decision tree performance on train and test sets with different tree_
↳ depths
from sklearn.datasets import make_classification
from sklearn.model_selection import train_test_split
from sklearn.metrics import accuracy_score
from sklearn.tree import DecisionTreeClassifier
from matplotlib import pyplot

```

```
[24]: # synthetic classification dataset
from sklearn.datasets import make_classification
# define dataset
X, y = make_classification(n_samples=10000, n_features=20, n_informative=5,
    ↪n_redundant=15, random_state=1)
# summarize the dataset
print(X.shape, y.shape)
```

(10000, 20) (10000,)

Use the `train_test_split()` function and split the data into 70 percent for training a model and 30 percent for evaluating it

```
[25]: # split into train test sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3)
```

```
[26]: # define lists to collect scores
train_scores, test_scores = list(), list()
# define the tree depths to evaluate
values = [i for i in range(1, 21)]
```

```
[27]: # evaluate a decision tree for each depth
for i in values:
    # configure the model
    model = DecisionTreeClassifier(max_depth=i)
    # fit model on the training dataset
    model.fit(X_train, y_train)
    # evaluate on the train dataset
    train_yhat = model.predict(X_train)
    train_acc = accuracy_score(y_train, train_yhat)
    train_scores.append(train_acc)
    # evaluate on the test dataset
    test_yhat = model.predict(X_test)
    test_acc = accuracy_score(y_test, test_yhat)
    test_scores.append(test_acc)
    # summarize progress
    print('>%d, train: %.3f, test: %.3f' % (i, train_acc, test_acc))
```

```
>1, train: 0.767, test: 0.769
>2, train: 0.808, test: 0.805
>3, train: 0.882, test: 0.881
>4, train: 0.903, test: 0.902
>5, train: 0.910, test: 0.905
>6, train: 0.913, test: 0.906
>7, train: 0.930, test: 0.922
>8, train: 0.945, test: 0.925
>9, train: 0.958, test: 0.928
>10, train: 0.964, test: 0.928
```

```

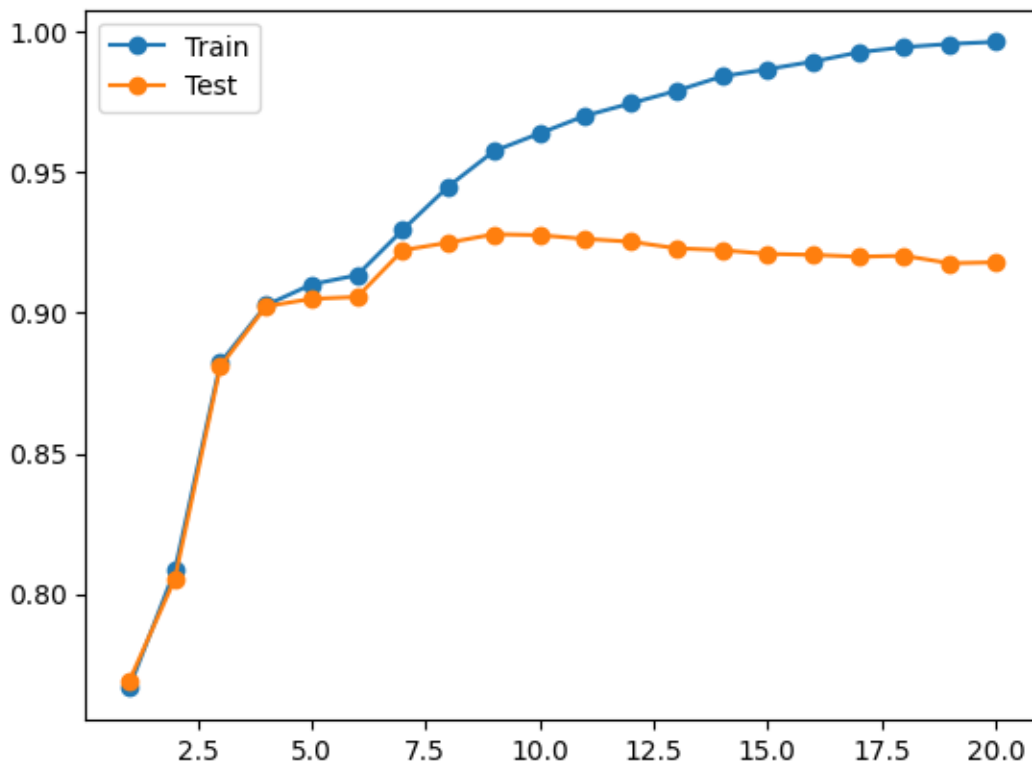
>11, train: 0.970, test: 0.926
>12, train: 0.975, test: 0.925
>13, train: 0.979, test: 0.923
>14, train: 0.984, test: 0.922
>15, train: 0.987, test: 0.921
>16, train: 0.989, test: 0.921
>17, train: 0.993, test: 0.920
>18, train: 0.995, test: 0.920
>19, train: 0.996, test: 0.918
>20, train: 0.996, test: 0.918

```

```

[28]: # plot of train and test scores vs tree depth
pyplot.plot(values, train_scores, '-o', label='Train')
pyplot.plot(values, test_scores, '-o', label='Test')
pyplot.legend()
pyplot.show()

```



```

[29]: from sklearn.tree import DecisionTreeClassifier

# Train a deep decision tree (Overfitting case)
deep_tree = DecisionTreeClassifier(max_depth=None) # No depth limit
deep_tree.fit(X_train, y_train)

```

```

print("Deep Tree Accuracy:", deep_tree.score(X_test, y_test))

# Train a pruned decision tree (Less overfitting)
pruned_tree = DecisionTreeClassifier(max_depth=3) # Limited depth
pruned_tree.fit(X_train, y_train)
print("Pruned Tree Accuracy:", pruned_tree.score(X_test, y_test))

```

Deep Tree Accuracy: 0.9143333333333333

Pruned Tree Accuracy: 0.881

car prediction data set - decision regressor

```
[30]: print(data.columns)
```

```

Index(['sepal_length', 'sepal_width', 'petal_length', 'petal_width', 'type'],
      dtype='object')

```

```
[31]: print(data.head()) # Check dataset preview
```

	sepal_length	sepal_width	petal_length	petal_width	type
1	5.1	3.5	1.4	0.2	Iris-setosa
2	4.9	3.0	1.4	0.2	Iris-setosa
3	4.7	3.2	1.3	0.2	Iris-setosa
4	4.6	3.1	1.5	0.2	Iris-setosa
5	5.0	3.6	1.4	0.2	Iris-setosa

```

[38]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import LabelEncoder
from sklearn.tree import DecisionTreeRegressor, plot_tree
from sklearn.metrics import mean_absolute_error, mean_squared_error, r2_score

# Load dataset & preprocess
data = pd.read_csv("carprediction.csv")
data.columns = data.columns.str.strip()
X, y = data.drop(columns=['MSRP']), data['MSRP']

# Encode categorical variables
for col in X.select_dtypes(include=['object']).columns:
    X[col] = LabelEncoder().fit_transform(X[col])

# Train-test split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
                                                    random_state=42)

# Train & Predict

```

```

model = DecisionTreeRegressor(max_depth=5, random_state=42)
model.fit(X_train, y_train)
y_pred = model.predict(X_test)

# Evaluation
print(f"MAE: {mean_absolute_error(y_test, y_pred)}")
print(f"RMSE: {np.sqrt(mean_squared_error(y_test, y_pred))}")
print(f"R2 Score: {r2_score(y_test, y_pred)}")

# Plot Decision Tree
plt.figure(figsize=(50, 10))
plot_tree(model, feature_names=X.columns, filled=True, fontsize=6)
plt.show()

```

MAE: 8495.966239509191
RMSE: 18655.36340215665
R² Score: 0.8539910358075516

